

A Case Study for Moving Targets

A Case Study for Moving Targets

Let's consider a more practical use case for MT

We will still tackle a synthetic problem, but one closer to practice

- In particular, given a classification problem
- We will require to have roughly balance class predictions

$$\left| \sum_{i=1}^m z_{ij} - \frac{m}{n_c} \right| \leq \beta \frac{m}{n_c}, \quad \forall j \in 1..n_c$$

- Where $z_{ij} = 1$ iff the classifier predicts class j for example i
- I.e. the result of an argmax applied to the output of a probabilistic classifier

...Basically, this the "very bad example" from the previous section

The Dataset

We will use the "wine quality" dataset from UCI

```
In [2]: data = util.load_classification_dataset(f'{data_folder}/winequality-white.csv', onehot_inputs=['  
display(data.head())
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality_3	quality_4	quality_5	qual
0	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.0010	3.00	0.45	8.8	0	0	0	1
1	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.9940	3.30	0.49	9.5	0	0	0	1
2	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.9951	3.26	0.44	10.1	0	0	0	1
3	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9	0	0	0	1
4	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9	0	0	0	1

- We will learn a model to predict wine quality
- There are 7 possible classes, represented via a one-hot encoding
- An ordinal encoding would be better, but our choice makes for a better example

The Dataset

We perform pre-processing as usual

```
In [3]: dtout = [c for c in data.columns if c.startswith('quality_')]
dtin = [c for c in data.columns if c not in dtout]
trl, tsl, scalers = util.split_datasets([data], fraction=0.7, seed=42, standardize=dtin)
tr, ts, scaler = trl[0], tsl[0], scalers[0]
tr.describe()
```

Out [3]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density
count	1.469000e+03	1.469000e+03	1.469000e+03	1.469000e+03	1.469000e+03	1.469000e+03	1.469000e+03	1.469000e+03
mean	5.235960e-16	1.475259e-16	-1.934766e-16	1.644551e-16	-1.015752e-16	-1.813843e-16	1.463167e-16	-8.416231e-15
std	1.000341e+00	1.000341e+00	1.000341e+00	1.000341e+00	1.000341e+00	1.000341e+00	1.000341e+00	1.000341e+00
min	-3.513954e+00	-1.981296e+00	-2.750607e+00	-1.170902e+00	-1.436585e+00	-1.916670e+00	-2.969922e+00	-2.338980e+00
25%	-6.396233e-01	-6.660564e-01	-5.421591e-01	-9.256225e-01	-4.519909e-01	-6.550002e-01	-7.262526e-01	-7.862395e-01
50%	-4.080447e-02	-1.601951e-01	-1.331873e-01	-2.306630e-01	-1.387109e-01	-8.151391e-02	-1.309932e-01	-5.839245e-02
75%	5.580144e-01	4.468384e-01	4.393732e-01	7.197965e-01	1.745692e-01	5.493210e-01	6.474229e-01	7.387734e-01
max	8.821714e+00	4.898418e+00	5.347035e+00	4.031074e+00	1.006527e+01	1.454239e+01	6.897646e+00	3.112941e+00

Dataset Balance

We can use the (avg. of) our constraint metric to assess the dataset balance:

$$\frac{1}{n_c} \sum_{j=1}^{n_c} \left| \sum_{i=1}^m \hat{z}_{ij} - \frac{m}{n_c} \right|, \quad \forall j \in 1..n_c$$

- Where \hat{z} are the class columns (one-hot encoding)

```
In [4]: bal_thr = 0.2
tr_true = np.argmax(tr[dtout].values, axis=1)
ts_true = np.argmax(ts[dtout].values, axis=1)
tr_bal_src = util.avg_bal_deviation(tr_true, bal_thr, nclasses=len(dtout))
ts_bal_src = util.avg_bal_deviation(ts_true, bal_thr, nclasses=len(dtout))
print(f'Original avg deviation: {tr_bal_src*100:.2f}% (training), {ts_bal_src*100:.2f}% (test)')
```

Original avg deviation: 101.42% (training), 98.71% (test)

- Our goal will be to push the balance deviation down to 20%
- I.e. we will assume $\beta = 0.2$ in the constraint

The Learner

Our "learner" will be a multilayer perceptron

The code can be found as usual in the `util` module

```
class MLP_Learner(object):  
    def __init__(self, hidden, epochs=20, batch_size=32,  
                  epochs_fine_tuning=None, verbose=0): ...  
  
    def fit(self, X, y): ...  
  
    def predict_proba(self, X): ...  
  
    def predict(self, X): ...
```


- We are using a standard scikit-learn API
- ...With the ability to use different #epochs for the first and subsequent training
- ...Which will prove useful later

The Learner

Let's start by checking how regular training fares

```
In [5]: hidden = [8, 8]
learner = util.MLPLearner(hidden=hidden, epochs=150)
learner.fit(tr[dtin].values, tr[dtout].values)
util.plot_training_history(learner.history, figsize=figsize)
```

2022-07-03 17:52:29.697366: I tensorflow/core/platform/cpu_feature_guard.cc:151] This TensorFlow binary is optimized with oneAPI Deep Neural Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX2 FMA
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.



The Learner

Now we can check the model performance

...In terms of both accuracy and degree of constraint violation

```
In [6]: tr_pred_prob = learner.predict_proba(tr[dtin])
        ts_pred_prob = learner.predict_proba(ts[dtin])
        tr_acc, tr_bal = util.mt_balance_stats(tr[dtout].values, tr_pred_prob, bal_thr)
        ts_acc, ts_bal = util.mt_balance_stats(ts[dtout].values, ts_pred_prob, bal_thr)
        print(f'Accuracy: {tr_acc:.2f} (training), {ts_acc:.2f} (test)')
        print(f'Classifier avg deviation: {tr_bal*100:.0f}% (training), {ts_bal*100:.0f}% (test)')
        print(f'Balance violation threshold: {bal_thr*100:.0f}%')
```

```
Accuracy: 0.59 (training), 0.53 (test)
Classifier avg deviation: 142% (training), 141% (test)
Balance violation threshold: 20%
```

- The accuracy is slightly above 50%
- ...But the balance violation is far larger than our threshold

The Master

The master step is implemented via an Or-tools model

```
def mt_balance_master(y_true, y_pred, bal_thr, alpha=1, time_limit=None, mode=...):  
    # Build a model  
    slv = pywraplp.Solver.CreateSolver('CBC')  
  
    ...  
    # Solve  
    status = slv.Solve()  
  
    ...  
    # Return the solution and stats  
    return sol, stats
```

- y_{true} corresponds to \hat{y} and y_{pred} to the current prediction vector
- The balance threshold bal_thr is a fractional value
- A mode parameter allows one to adjust a bit the problem behavior

The Master

The master step is implemented via an Or-tools model

```
def mt_balance_master(y_true, y_pred, bal_thr, alpha=1, time_limit=None, mode='gradient')
    ...
    # Build target variables
    z = {(i,j) : slv.IntVar(0, 1, f'z[{i},{j}]')} for i in range(ns) for j in range(nc)
    # Unique class constraints
    for i in range(ns):
        slv.Add(sum(z[i, j] for j in range(nc)) == 1)
    ...
```

- We are using integer variables for the targets
- This is partly for sake of simplicity and scalability
- ...And partly since it is compatible with DT models as well
- Using continuous (probabilistic) targets is possible, but harder

The Master

The master step is implemented via an Or-tools model

```
def mt_balance_master(y_true, y_pred, bal_thr, alpha=1, time_limit=None, mode='gradient')
    ...
    # Add the balance constraint
    ref = ns / nc
    for j in range(nc):
        slv.Add(sum(z[i, j] for i in range(ns)) <= ref + ref * bal_thr)
        slv.Add(sum(z[i, j] for i in range(ns)) >= ref - ref * bal_thr)
    ...
```

- The balance constraint is implemented via two inequalities
- Class counts are easy to compute by relying on integer one-hot targets

The Master

The master step is implemented via an Or-tools model

```
def mt_balance_master(y_true, y_pred, bal_thr, alpha=1, time_limit=None, mode='gradient')
    ...
    # Build the gradient-based part of the objective
    loss_t = 0
    for i in range(ns):
        for j in range(nc):
            loss_t += 0.5 * dsgn(y_pred[i, j], y_true[i, j]) * (z[i, j] - y_pred[i, j])
    ...
```

- As a loss L for the master, we use $\|y - \hat{y}\|$
- ...Hence the gradient is given by the sign of the difference $y - \hat{y}$
- If $y_i = \hat{y}_i$ for some i , the `dsgn` function has a customized behavior

The Master

The master step is implemented via an Or-tools model

```
def mt_balance_master(y_true, y_pred, bal_thr, alpha=1, time_limit=None, mode='gradient')
    ...
    # Build the quadratic part of the objective
    loss_p = 0
    for i in range(ns):
        for j in range(nc):
            loss_p += z[i, j] * (1-y_pred[i, j])**2 + (1-z[i, j]) * (0-y_pred[i, j])**2
    # Define the cost function
    slv.Minimize(alpha * loss_t + loss_p)
    ...
```

- The quadratic part of the objective is easy to linearize
- ...Since we are using integer z variables

Loss-driven Projection

A simpler approach to inject constraints in the ML model...

...Starts by directly "projecting" the ground truth \hat{y} in feasible space

- The projection can be done using the loss itself as a distance:

$$\operatorname{argmin}_z \{ L(z, \hat{y}) \mid z \in C \}$$

By doing this, we can obtain the **best possible feasible target vector**

```
In [7]: zp, stats = util.mt_balance_master(tr[dtout].values, tr[dtout].values, bal_thr, time_limit=10, n
tmp_acc, tmp_bal = util.mt_balance_stats(tr[dtout].values, zp, bal_thr)
print(f'Accuracy: {tmp_acc:.2f}, Balance deviation: {tmp_bal*100:.2f}%, Optimal solution: {stats
```

```
Accuracy: 0.58, Balance deviation: 16.80%, Optimal solution: False
```

- In our case, the solution is not optimal (there is a time limit)
- ...But it's very close to optimality

Loss-driven Projection

Then, the simple approach consists training against this "ideal" vector

The method is implemented in `util` as part of the MT code:

```
In [8]: learner_prj = util.MLPLearner(hidden=hidden, epochs=600)
util.mt_balance(tr[dtin].values, tr[dtout].values, learner_prj, bal_thr, mode='projection', mast
tr_pred_prob = learner_prj.predict_proba(tr[dtin])
ts_pred_prob = learner_prj.predict_proba(ts[dtin])
tr_acc, tr_bal = util.mt_balance_stats(tr[dtout].values, tr_pred_prob, bal_thr)
ts_acc, ts_bal = util.mt_balance_stats(ts[dtout].values, ts_pred_prob, bal_thr)
print(f'Accuracy: {tr_acc:.2f} (training), {ts_acc:.2f} (test)')
print(f'Classifier balance violation: {tr_bal*100:.0f}% (training), {ts_bal*100:.0f}% (test)')
```

Accuracy: 0.40 (training), 0.40 (test)

Classifier balance violation: 99% (training), 100% (test)

- The accuracy is lower, since we have confounded the input/output relation
- ...And the violation is still not fine
 - The ML model bias prevents it from reaching the adjusted target
 - ...But too little bias may result in overfitting

Moving Targets

Let's now test the actual MT method

We use fewer training epochs after the first `fit` call to speed up the process

```
In [11]: alpha, max_iter = 5, 10
epochs, epochs_fine_tuning = 150, 30
learner_mt = util.MLPLearner(hidden=hidden, epochs=epochs, epochs_fine_tuning=epochs_fine_tuning,
util.mt_balance(tr[dtin].values, tr[dtout].values, learner_mt, bal_thr, alpha,
               max_iter, master_tlim=2, verbose=1);
```

```
(#1) l-acc: 0.35, l-bal: 0.43, l-time: 2.30s, m-acc: 0.58, l-m dist: 3.00, m-time: 2.42s
(#2) l-acc: 0.32, l-bal: 0.21, l-time: 1.96s, m-acc: 0.58, l-m dist: 3.07, m-time: 2.64s
(#3) l-acc: 0.31, l-bal: 0.17, l-time: 1.96s, m-acc: 0.57, l-m dist: 2.96, m-time: 2.39s
(#4) l-acc: 0.31, l-bal: 0.17, l-time: 1.89s, m-acc: 0.54, l-m dist: 2.96, m-time: 2.71s
(#5) l-acc: 0.31, l-bal: 0.16, l-time: 1.89s, m-acc: 0.49, l-m dist: 2.94, m-time: 2.69s
(#6) l-acc: 0.33, l-bal: 0.17, l-time: 2.04s, m-acc: 0.44, l-m dist: 3.02, m-time: 2.52s
(#7) l-acc: 0.34, l-bal: 0.17, l-time: 2.04s, m-acc: 0.39, l-m dist: 3.06, m-time: 2.75s
(#8) l-acc: 0.34, l-bal: 0.18, l-time: 1.87s, m-acc: 0.37, l-m dist: 3.11, m-time: 2.37s
(#9) l-acc: 0.34, l-bal: 0.17, l-time: 1.91s, m-acc: 0.36, l-m dist: 3.12, m-time: 2.68s
(#10) l-acc: 0.34, l-bal: 0.17, l-time: 1.03s, m-acc: 0.35, l-m dist: 3.12, m-time: 2.31s
```

- Constraint satisfaction improves across iterations
- The learner/master accuracy tends to decrease/increase

Moving Targets

Let's check generalization over the test set

```
In [12]: tr_pred_prob = learner_mt.predict_proba(tr[dtin])
         ts_pred_prob = learner_mt.predict_proba(ts[dtin])
         tr_acc, tr_bal = util.mt_balance_stats(tr[dtout].values, tr_pred_prob, bal_thr)
         ts_acc, ts_bal = util.mt_balance_stats(ts[dtout].values, ts_pred_prob, bal_thr)
         print(f'Accuracy: {tr_acc:.2f} (training), {ts_acc:.2f} (test)')
         print(f'Classifier balance deviation: {tr_bal*100:.0f}% (training), {ts_bal*100:.0f}% (test)')
```

Accuracy: 0.34 (training), 0.25 (test)
Classifier balance deviation: 17% (training), 9% (test)

We do have some overfitting in terms of accuracy

- This is due to the fact that we have very restrictive constraints
- ...And they directly oppose information in the data

Constraint satisfaction generalizes without issues

- ...And this is a big deal!

Some References

- Boyd, Stephen, et al. "Distributed optimization and statistical learning via the alternating direction method of multipliers." Foundations and Trends® in Machine learning 3.1 (2011): 1-122.
- Fabrizio Detassis, Michele Lombardi, Michela Milano: Teaching the Old Dog New Tricks: Supervised Learning with Constraints. AAAI 2021: 3742-3749
- Aghaei, S.; Azizi, M. J.; and Vayanos, P. 2019. Learning optimal and fair decision trees for non-discriminative decision-making. In Proceedings of the AAAI Conference on Artificial Intelligence, volume 33, 1418–1426.
- Kamiran, F.; and Calders, T. 2012. Data preprocessing techniques for classification without discrimination. Knowledge and Information Systems 33(1): 1–33